



# Vault Self-Managed Getting Started & Technical Overview

October 2021

*Copyright © 2021 HashiCorp*

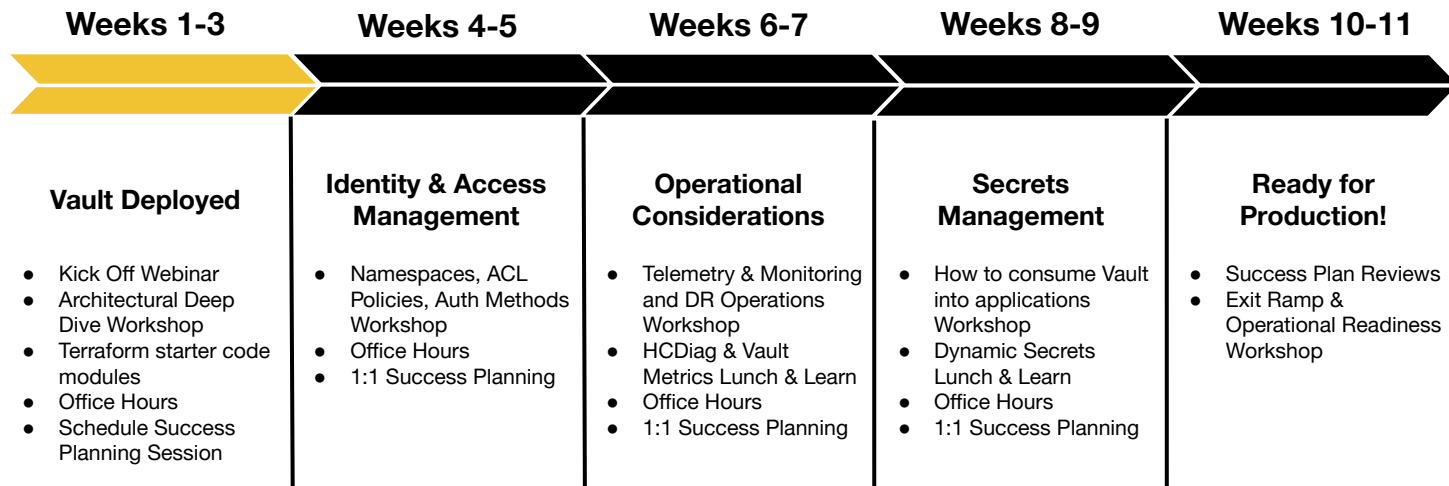


---

# Agenda

1. Overview
2. Architecture
3. Deployment Patterns
4. Operations
5. Next Steps

# Vault Enterprise Path to Production



01



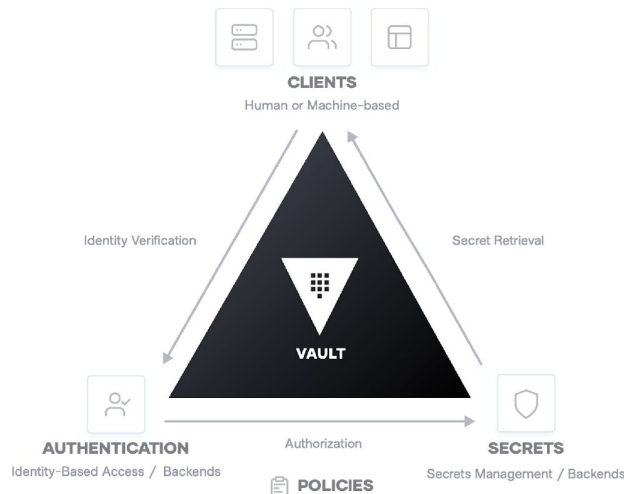
# Vault Overview

# Overview

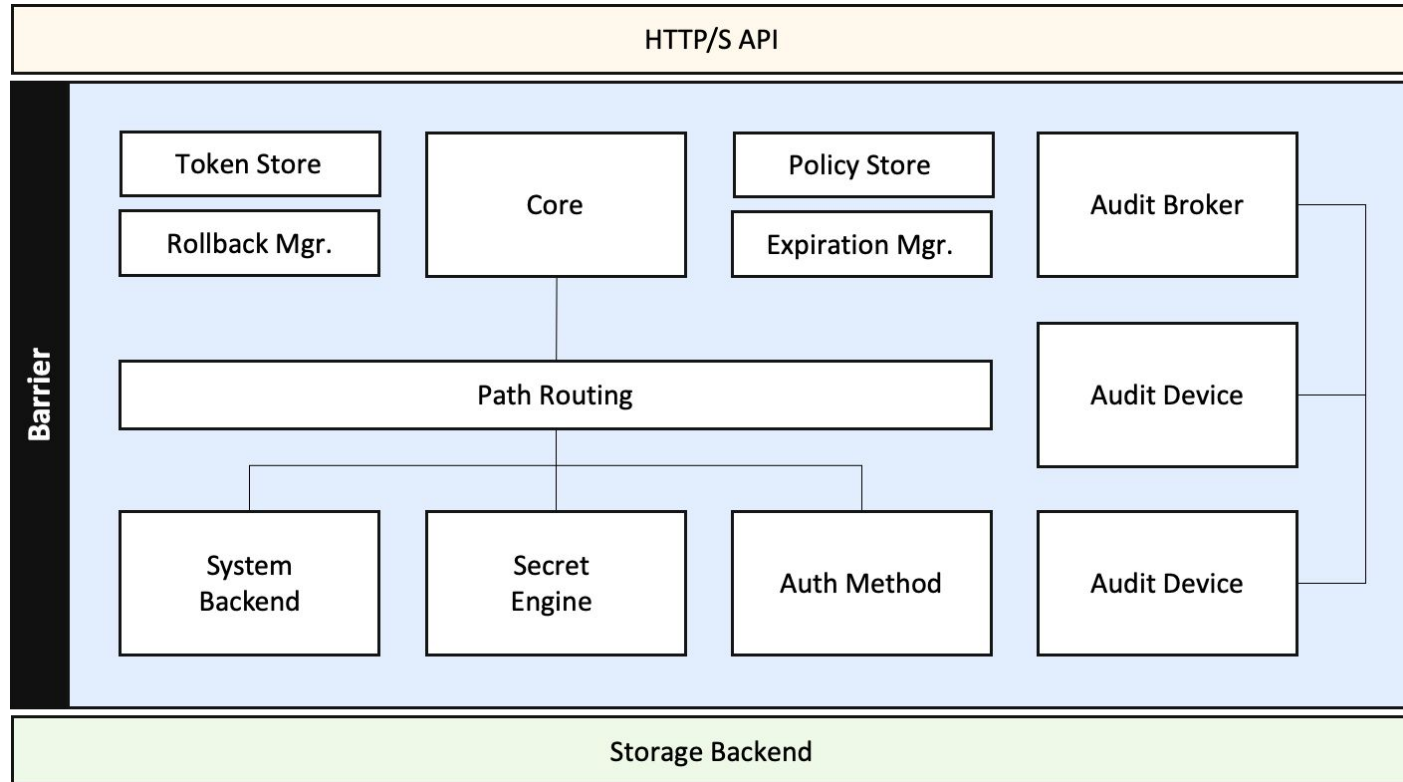


Vault tightly controls access to secrets and encryption keys by authenticating against trusted sources of identity such as Active Directory, LDAP, and cloud identity platforms.

Vault enables fine grained authorization of which users and applications are permitted access to secrets and keys.



# Architecture & Cryptographic Barrier



# Vault Security Model



- It's all about access to the **Encryption Key**.
- Configuring `cap_ipc_lock=+ep`, `LimitNOFILE`, and `LimitMEMLOCK` prevent Memory Swapping to Disk, so secrets are not written in plain text to disk.
- The Vault Encryption Key is stored in memory in **PLAIN TEXT**
  - This is done for performance
  - Root access to an unlocked vault server could compromise this
  - Isolation technologies which allow reading of memory could compromise this (VM issues, but principally Kubernetes)
- Master Root Key protects the Encryption key, so it also must be secure
- Auto-Unseal shifts risk profile

# Cryptography Security Model



- **Vault uses publicly available cryptographic technologies**
- P vs NP - Good cryptographic algorithms are exponential in difficulty to solve but polynomial in difficulty to validate answers for.
- Numerous algorithms (SHA1) were exposed to have defects that allowed them, or a subset of them to be reduced to polynomial difficulty problems
- Short encryption keys and faster computers has made brute-forcing older encryption standards possible.
- Software based random number generations suffer from a lack of randomness.



02

# Vault Architecture

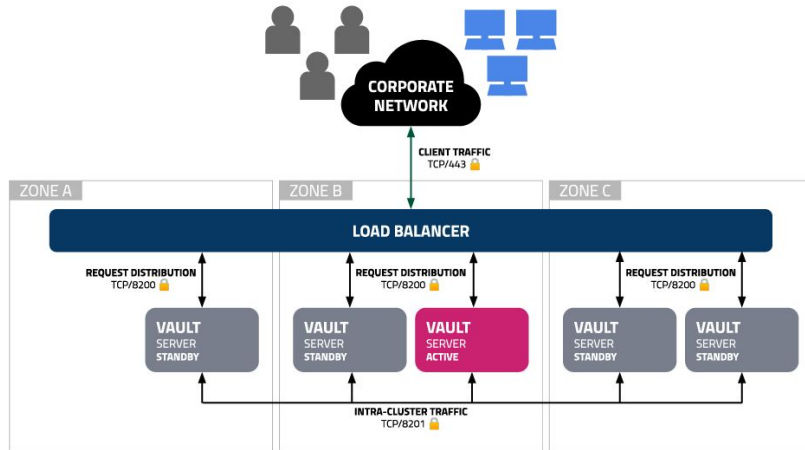
# Vault Cluster Architecture



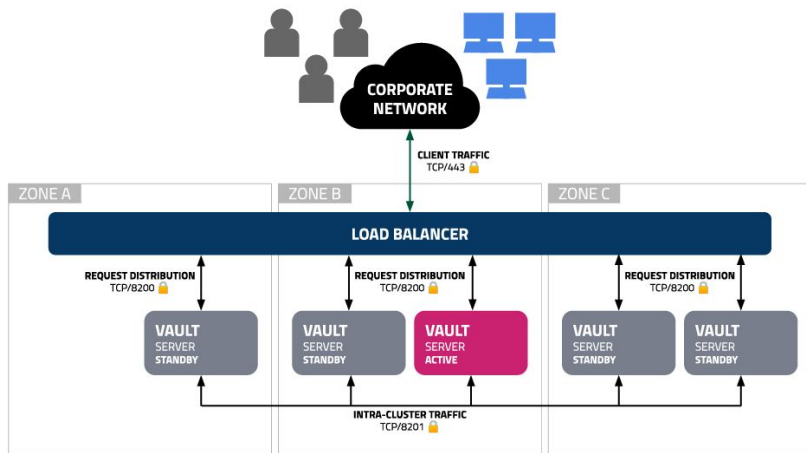
## Integrated Storage Reference Architecture

This architecture provides a highly resilient and scalable Vault deployment.

- Able to withstand loss of two nodes or an entire Availability Zone (AZ).
- Integrated Storage requires high-bandwidth, low latency (<8ms RT) connection.



# Vault Integrated Storage Architecture



## 5 Vault Servers, 3 Availability Zones

### On-premise Considerations

A single region in a On-premise Deployment is considered a geographically separate collection of one or more Availability Zones.

- An AZ would be considered a single failure domain containing it's own compute, storage, network, power, etc.
- A high bandwidth, low latency (<8ms RT) connection should exist between each AZ.



# Sizing

Per instance sizing  
recommendations

	Small (Dev/Test/Staging /QA)	Large (Production)
<b>CPU</b>	2 - 4 Core	4 - 8 Core
<b>Memory</b>	8 - 16 GB RAM	32 - 64 GB RAM
<b>Disk Capacity</b>	100+ GB	200+ GB
<b>Disk IO</b>	3000+ IOPS	10000+ IOPS
<b>Disk Throughput</b>	75+ MB/s	250+ MB/s

# Cloud Instance Sizing



Provider	Size	Instance/VM Types	Disk Volume Specs
AWS	Small	m5.large, m5.xlarge	100+ GB gp3, 3000 IOPS, 125 MB/s
	Large	m5.2xlarge, m5.4xlarge	200+ GB gp3, 10000 IOPS, 250 MB/s
Azure	Small	standard_d2s_v3, standard_d4s_v3	1024 GB Premium_LRS
	Large	standard_d8s_v3, standard_d16s_v3	1024 GB Premium_LRS
GCP	Small	n2-standard-2, n2-standard-4	500 GB pd-balanced
	Large	n2-standard-8, n2-standard-16	1000 GB pd-ssd

# Performance Considerations



## Profile Workloads

As you scale the adoption of Vault throughout your organization, you will have varying workloads access Vault. Telemetry monitoring should be leveraged to ensure proactive monitoring of Vault Cluster resources. Additionally, as you onboard new applications/services/teams/users to Vault, take time to profile the usage patterns to ensure optimal authentication and consumption patterns are used.

## External Systems

Depending on the Authentication Methods and Secrets Engines used by your organization, you will likely have dependency on other external systems for Vault requests to be completed. Ensure telemetry is enabled on those services and proactively monitor for performance issues.

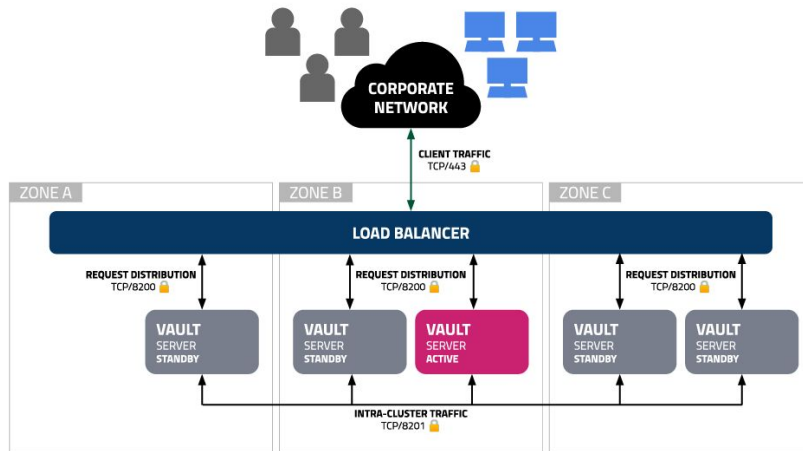
# Networking



## Networking Considerations

Integrated Storage can enable you to simplify your networking requirements compared to other storage backends.

- <8ms RT network connection required to ensure Raft Storage remains consistent across all Vault Nodes.
- Restrict communication to only required ports and CIDRs
- Standard HTTPS TLS encryption should be used to protect network traffic



# Networking Requirements



Source	Destination	Port	Protocol	Direction	Purpose
Client Machines	Load Balancer	443	tcp	incoming	Request distribution
Load Balancer	Vault Servers	8200	tcp	incoming	Vault API
Vault Servers	Vault Servers	8200	tcp	bidirectional	Cluster Bootstrapping
Vault Servers	Vault Servers	8201	tcp	bidirectional	Raft, replication, request forwarding
Vault Servers	External Systems	various	various	various	External APIs



# High Availability



## HA Strategy for Client Access

Vault does not include built in load balancing capabilities. To ensure Vault can meet the highest levels of reliability and stability, either an external load balancer or Consul should be used to distribute client requests.

- TLS should terminate at Vault and not the load balancer to ensure end-to-end encryption
- Use Vault's health endpoint to determine active node and node health  
<https://<vaultnode>:8200/v1/sys/health>



# Scaling Considerations

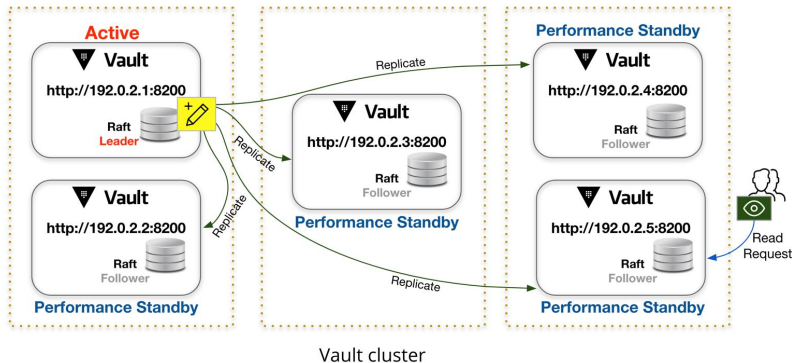


Managed scaling services should be leveraged when deploying in a cloud environment to ensure the Vault cluster remains populated with health nodes.

- Additional nodes will not increase performance. Scaling services will increase your resiliency but not improve performance.
- Do not replace all instances at once in a scaling group otherwise you will have to restore from a snapshot

Cloud	Managed Auto Scaling Service
<b>AWS</b>	Auto Scaling Group
<b>Azure</b>	Virtual Machine Scale Sets
<b>GCP</b>	Managed Instance Groups

# Scaling Performance Standby Nodes



## Horizontal scalability for read requests

Performance Standby Nodes can be used to respond to read-only requests. Performance Standby Nodes are enabled by default and process read-only requests locally. Write requests that modify shared state in Vault will be forwarded to the singular Active Node.

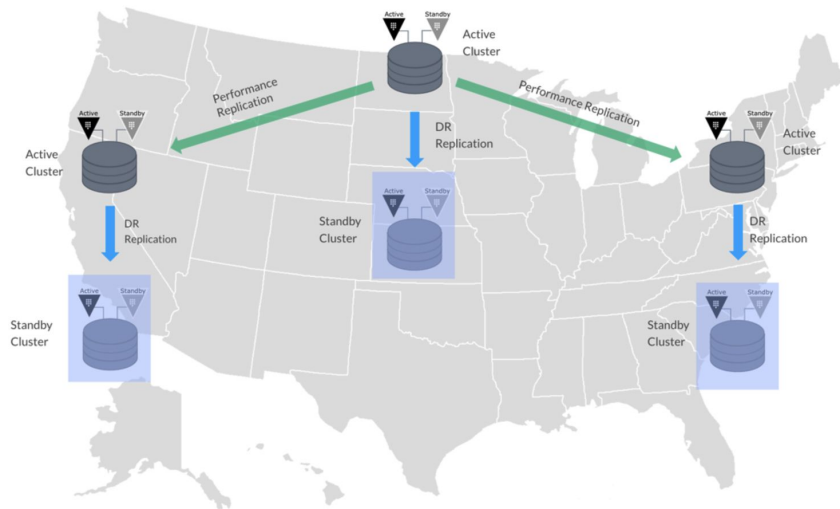
- Integrated Storage uses eventual consistency and data may not be available across all nodes immediately
- Vault 1.7+ includes multiple methods to control how requests are handled

# Multiple Regions



## Vault Replication

Vault can be extended to multiple regions using replication. The primary cluster uses asynchronous replication to ship data to the secondaries. Multiple replication modes can be combined to provide resilience and performance.



# Replication Types



## Disaster Recovery

Provides a warm standby cluster that contains all data from the primary Vault cluster. **It is strongly recommended to deploy at least one DR cluster.**


## Performance Replication

Provides an active Vault cluster with shared state of the primary. This includes secrets, authentication methods, policies, and other shared data. Token and leases are not replicated to performance secondaries.




# Replication

## Determine Replication Mode



### DR Replication

- RPO/RTO requirements
- Tier 0 or 1 apps be accessing secrets from Vault



### Performance Replication

- Source of Requests
- Performance requirements for request and response
- Types of workloads

03

# Deployment Patterns

# Recommended Patterns



## **Immutable Builds**

Tooling like Packer can be used to build immutable images of Vault and perform blue/green deployment using your existing CI/CD orchestration. This can streamline your lifecycle processes. However, when using integrated storage, you will need to take measures to ensure quorum is maintained as new image versions are introduced to the cluster.

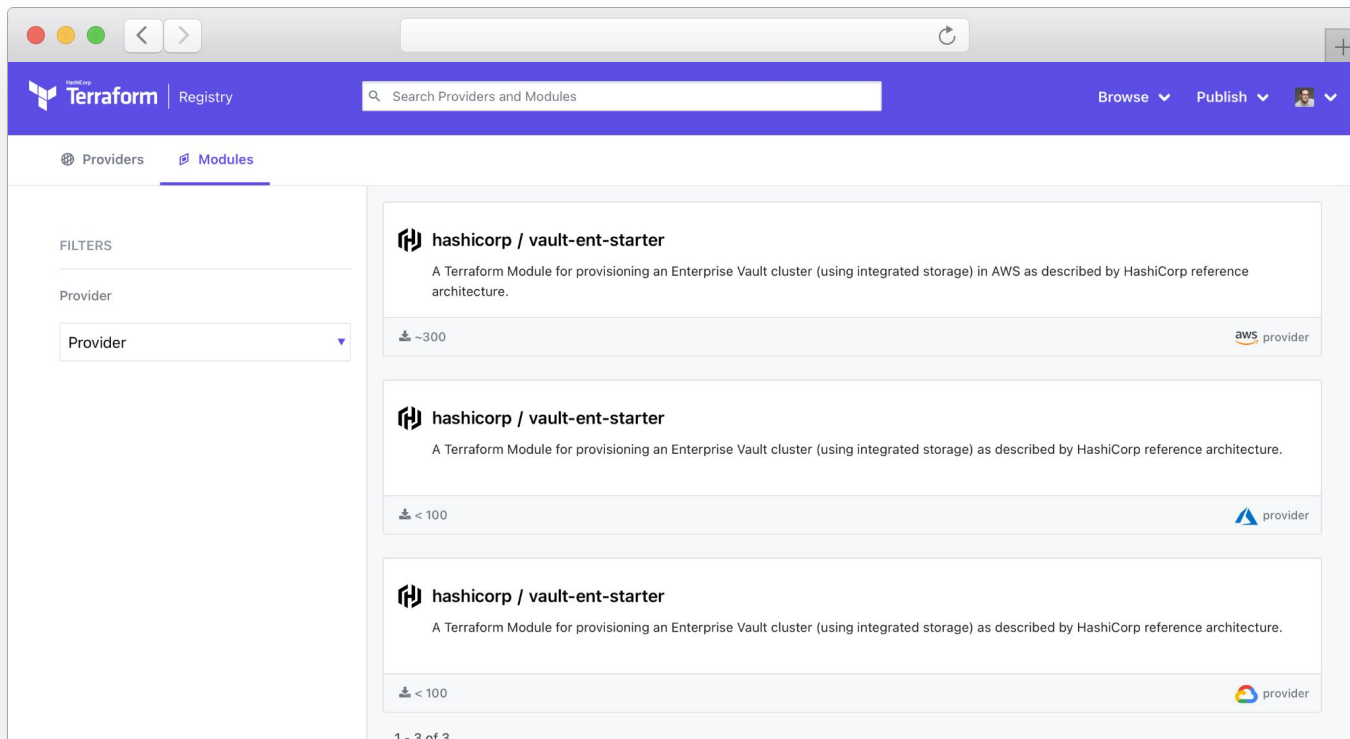
## **Configuration Management**

For organizations who have not adopted the above pattern, Vault can be integrated into your configuration management patterns to install, upgrade, and configure Vault.



# Terraform Modules

Quickly deploy Vault cluster based on reference architecture



# Vault Helm Chart

Deploy Vault Reference Architecture inside Kubernetes



The screenshot shows the GitHub repository for the Vault Helm Chart, maintained by HashiCorp. The repository is public and has 50 watches, 622 stars, and 515 forks. It features 98 issues, 35 pull requests, and 25 tags. The main branch is 'main'. The repository contains a file tree with directories like .circleci, .github, templates, and test, and files like .gitignore, .helmignore, CHANGELOG.md, CONTRIBUTING.md, and Chart.yaml. The 'About' section describes the chart as a Helm chart to install Vault and other associated components, with a Readme and MPL-2.0 License. The 'Releases' section shows the latest release is v0.16.1, published 9 days ago, with 24 releases in total. The 'Packages' section indicates no packages are published.

hashicorp / vault-helm Public

Watch 50 Star 622 Fork 515

Code Issues 98 Pull requests 35 Actions Projects Security Insights

main 4 branches 25 tags

Code

About

Helm chart to install Vault and other associated components.

Readme

MPL-2.0 License

Releases 24

v0.16.1 Latest 9 days ago

+ 23 releases

Packages

No packages published

File	Commit	Time
.circleci	fix chart publish job (#620)	9 days ago
.github	Update jira sync github action (#411)	11 months ago
templates	Adding support for the old leader-elect (607)	23 days ago
test	vault-helm 0.16.1 release (#619)	9 days ago
.gitignore	feature: Support configuring various properties as YAML directly. (#5...	3 months ago
.helmignore	Ignore bin dirs	3 years ago
CHANGELOG.md	vault-helm 0.16.1 release (#619)	9 days ago
CONTRIBUTING.md	vault-helm default branch is now main (#618)	11 days ago
Chart.yaml	vault-helm 0.16.1 release (#619)	9 days ago



---

# Upgrades

Major releases of Vault are released quarterly and we release minor releases monthly and as needed

The process to update Vault should be automated as much as possible to ensure Vault remain patched with the latest bug and security fixes. We ensure upgrade compatibility with N-2 major releases.

Prior to performing any upgrade it is recommended to review the changelog and version specific upgrade guide, test against version in QA environment, and ensure up to date snapshots are available in case you need to restore.

04

# Operations

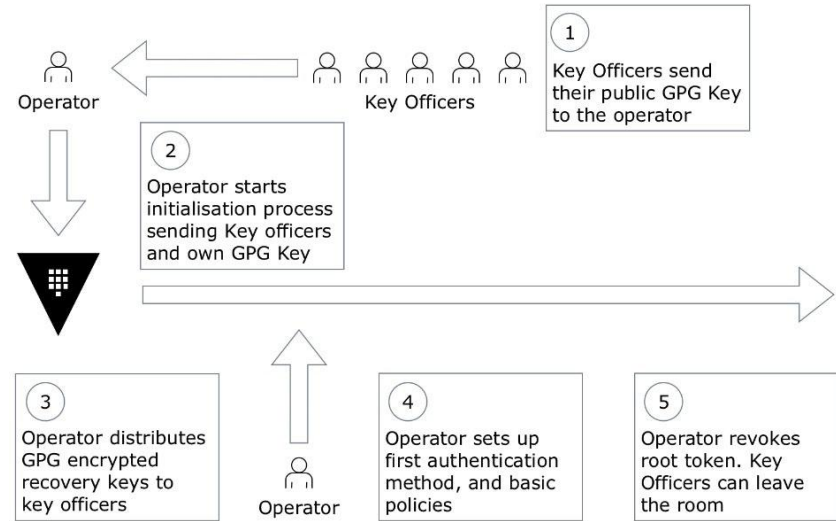
# Initialization Ceremony



## Vault Cluster Initialization Process

The initialization process only occurs once in a Vault installation. During this process, unseal or recovery keys and root token are returned to the operator. Keys are then distributed to key officers who are the guardians of Vault trust.

- Recommended to complete initialization process in single room with Vault operator and key holders until root token is revoked.
- Recovery Keys are returned when Vault is initialized while using Cloud KMS or HSM. Recovery keys should be protected the same as unseal keys.





# Root Token Generation

## TERMINAL

```
> vault operator init  
Unseal Key 1: Ly7wgNFzKVcw95nv6fLTQ/lsf49Wn4JaIEYGPm15pSzn  
Unseal Key 2: JWeteKjgpFXI2wY2I16j8JCCy92P04GxGCyXvLCoHp1L  
Unseal Key 3: zLkMb09Lcr3QRwIgw7KBPy5jRD9aUttt010HZ4dusvx  
Unseal Key 4: 0J5fD29c5ZisK11jL13K0XOmIWu66PfA6NBV3UEK7f/f  
Unseal Key 5: ahR01B203Kzxv0a0HgBLUDmByxhFdeVOVeA3l6PMIKMn
```



```
Initial Root Token: s.dZ1m130RBFkF0rQeWtLF3uiA
```

Vault initialized with 5 key shares and a key threshold of 3. Please securely distribute the key shares printed above. When the Vault is re-sealed, restarted, or stopped, you must supply at least 3 of these keys to unseal it before it can start servicing requests.

Vault does not store the generated master key. Without at least 3 key to reconstruct the master key, Vault will remain permanently sealed!

It is possible to generate new unseal keys, provided you have a quorum of existing unseal keys shares. See "vault operator rekey" for more information.



---

# Root Token Handling Practices

The root token is returned to the operator during the initialization ceremony. This token can do **anything** in Vault and its usage should be closely monitored.

- Once operator has configured a secondary authentication method and granted administrators sudo access, almost all operations can be performed.
- Best practice is **NOT** persisting the root token.
- Generate a root token only when absolutely necessary.



---

# Production Readiness

**Critical items to have in  
place before production  
go-live**

## **Backup**

Automated Integrated Storage Snapshots, included in your enterprise subscription, can take periodic snapshots of your data.

- Determine where snapshot files will be stored
- Configure based off your RPO/RTO requirements.
- If snapshot is restored, the unseal keys that were valid at the time of the snapshot will be used to unseal.





---

# Automated Integrated Storage Snapshots

```
TERMINAL

> vault write \
    sys/storage/raft/snapshot-auto/config/testsnap \
    storage_type=local \
    file_prefix=testsnappy \
    interval=120m \
    retain=7 \
    local_max_space=1000000 \
    path_prefix=/opt/vault/
```



---

# Production Readiness

**Critical items to have in place before production go-live.**

## Monitoring

Vault should be monitored closely to ensure the service remains healthy and available in production.

- Telemetry - Export telemetry data to solution that can analyze and identify trends overtime.
- Log Analytics - Capture app logs and system logs and perform analysis on the log files for useful signals.
- Active Health Checks - Query health endpoints to get the health of nodes and route traffic to active node.



---

# Production Readiness

**Critical items to have in place before production go-live.**



## **Auditing**


Vault sends audit information to a SIEM system or logging backend.




- Determine audit devices that will be used.
- Vault will not respond if the audit device is unavailable, use multiple audit devices to ensure Vault remains available.
- Sensitive fields are HMAC, Selectively determine if any HMAC fields need to be exposed






05

# Next Steps

HashiCorp | Discuss [Sign in](#)  

 Information on HashiCorp Cloud Platform (HCP) use cases, Q&A, and best practices discussions. Please note that offerings may be in various release lifecycles. Categorize your question or comment under [HCP Consul](#) or [HCP Vault](#) subcategories. If your question or comment relates to networking, access control, or billing, post in this category. For support requests, please [open a ticket](#).

HashiCorp Cloud Platform (HCP)  All  Tags  | Latest Top

Topic	Replies	Views	Activity
 <b>About the HashiCorp Cloud Platform (HCP) category</b> HashiCorp Cloud Platform (HCP) Information on HashiCorp Cloud Platform (HCP) use cases, Q&A, and best practices discussions. Please note that offerings may be in various release lifecycles. Categorize your question or comment under HCP Consul or HCP V... read more	 1	387	May 24
HCP Vault "per-client" pricing HCP Vault	 0	39	9d
<b>Failing to use HCP Consul as my terraform backend</b> HashiCorp Cloud Platform (HCP)	 1	72	12d
<b>Does HCP support Automation APIs in AWS</b> HashiCorp Cloud Platform (HCP) vault	 0	73	27d



# Discuss

Engage with the HashiCorp Cloud community including HashiCorp Architects and Engineers.


[discuss.hashicorp.com](https://discuss.hashicorp.com)

# Learn

Step-by-step guides to accelerate deployment of Vault



A screenshot of a web browser displaying the HashiCorp Learn website. The browser window has a title bar with standard macOS window controls (red, yellow, green buttons) and navigation arrows. The website header includes the HashiCorp logo, the text 'HashiCorp Learn', a 'Browse tutorials' dropdown menu, a search bar with a magnifying glass icon and the word 'Search', a document icon, and a blue 'Sign in' button. The main content area is divided into a left sidebar and a main article. The sidebar has a 'Vault' section with a sub-header 'GET STARTED' and links for 'CLI Quick Start', 'HCP Vault', and 'UI Quick Start'. Below this is a 'USE CASES' section with links for 'ADP', 'Data Encryption', and 'Secrets Management'. At the bottom of the sidebar is a 'CERTIFICATION PREP' section. The main article is titled 'Deploy Cluster with Integrated Storage' in a large, bold font. Below the title, it shows a clock icon with '2 HR 15 MIN' and a document icon with '12 TUTORIALS'. The article text begins with 'If you are responsible for setting up and maintaining a Vault cluster using integrated storage as a persistence layer, get started here.' Below the text is a card for a specific tutorial titled 'Vault with Integrated Storage Reference Architecture' with a '15 MIN' duration. The card text starts with 'This guide describes architectural best practices for implementing Vault using the Integrated Storage (Ref) storage backend.'



# Support

<https://support.hashicorp.com>

# Need Additional Help?



## Customer Success

Contact our Customer Success Management team with any questions. We will help coordinate the right resources for you to get your questions answered.

[customer.success@hashicorp.com](mailto:customer.success@hashicorp.com)

## Technical Support

Something not working quite right? Engage with HashiCorp Technical Support by opening a new ticket for your issue at [support.hashicorp.com](https://support.hashicorp.com).



The image features a dark, textured background with a pattern of small, light-colored dots. In the upper left corner, there are several parallel, diagonal gold lines that form a series of nested 'V' shapes. The text 'Q & A' is prominently displayed in the lower left area in a large, white, sans-serif font.

# Q & A



# Thank You

[hello@hashicorp.com](mailto:hello@hashicorp.com)

[www.hashicorp.com](http://www.hashicorp.com)